

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

a) Source code file names:

1. create\_data.py
2. HoG.py
3. neural\_net.py
4. run\_NN.py

b) Instructions on how to compile and run the code:

1. First place above four files in a folder.
2. Inside this folder create four folders as 'positive\_train', 'negative\_train', 'positive\_test', and 'negative\_test'.
3. Place all the positive train images in positive train folder and likewise place the other images in the respective folders.
4. First run the create\_data.py file. This file will create two csv files 'train.csv' and 'test.csv'.
5. Then run the 'run\_NN.py' file. It will display the accuracy and the predictions as outputs.
6. Or you can run the CV\_Project2.bat batch file which will do steps 4 and 5 for you.

c) Answers:

1. How did you initialize the weight values of the network?  
→ Random weights from normal distribution, 250 hidden neurons, and learning rate of 0.01
2. How many iterations (or epochs) through the training data did you perform  
→ After trying various combination, epoch was decided to be 400.
3. How did you decide when to stop training?  
→ When either number of iterations (epochs) are completed or the difference between successive errors is less than  $10^{-6}$
4. Based on the output value of the output neuron, how did you decide on how to classify the input image into human or not-human?  
→ The output neuron uses sigmoid as activation function so if the value is less than 0.5 the image is non-human else it is human.

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

d) Classification results on test images:


Test Image	Output value	Classification
crop_000010b	0.97051742	1
crop001008b	0.98690437	1
crop001028a	0.78469989	1
crop001045b	0.88945977	1
crop001047b	0.91908398	1
00000053a_cut	0.41748846	0
00000062a_cut	0.22092658	0
00000093a_cut	0.03816118	0
no_person__no_bike_213_cut	0.47318405	0
no_person__no_bike_247_cut	0.08421976	0

e) Comments:

Training the neural network may take some time depending upon the size of data and the terminating conditions.

Above is the best accuracy I got while trying different hyper parameters with random initial weights.

f) Normalized test images:



Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

g) Code:

1. HoG.py:

```
# importing required libraries
import numpy as np
from skimage.io import imread
from PIL import Image
import math
import sys

class HOG:
    def __init__(self, image_name):
        self.image_name = image_name

    # function to convert color image to greyscale
    def greyscale_operation(self, ip_img):
        op_img = np.zeros((ip_img.shape[0], ip_img.shape[1]))
        for i in range(ip_img.shape[0]):
            for j in range(ip_img.shape[1]):
                op_img[i][j] = np.round_(0.299*ip_img[i][j][0] + 0.587*ip_img[i][j][1] +
0.114*ip_img[i][j][2])
            return op_img

    # function for implementing gradient operation using prewitt's edge detector
    def gradient_operation(self, ip_img):
        # prewitt's vertical and horizontal kernels
        Gx_kernel = ([-1, 0, 1] * 3)
        Gy_kernel = ([1, 1, 1], [0, 0, 0], [-1, -1, -1])

        Gx = np.zeros((ip_img.shape[0], ip_img.shape[1]))
        Gy = np.zeros((ip_img.shape[0], ip_img.shape[1]))
        G = np.zeros((ip_img.shape[0], ip_img.shape[1]))
        Theta = np.zeros((ip_img.shape[0], ip_img.shape[1]))

        # calculating Gx and Gy using prewitt's edge detector
        for i in range(ip_img.shape[0]):
            for j in range(ip_img.shape[1]):
                # pixels for which part of the prewitt's mask goes outside of the image border
                if i < 1 or j < 1 or ip_img.shape[0] - i <= 1 or ip_img.shape[1] - j <= 1:
                    continue
                else:
                    arr1 = ip_img[i - 1:i + 2, j - 1:j + 2]
```

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

```
Gx[i][j] = np.sum(np.multiply(arr1, Gx_kernel))/3
Gy[i][j] = np.sum(np.multiply(arr1, Gy_kernel))/3
G[i][j] = round(math.sqrt(Gx[i][j]**2 + Gy[i][j]**2)/math.sqrt(2))
if Gx[i][j] == 0 and Gy[i][j] == 0:
    Theta[i][j] = 0
else:
    theta = np.arctan2(Gy[i][j], Gx[i][j])*180/np.pi
    if theta < 0:
        theta += 180
    if theta >= 170:
        theta -= 180
    Theta[i][j] = theta

return G, Theta

# HoG implementation
def hog_operation(self, ip_img, ip_theta):
    # HoG bins table
    bins_table = {0: [-10, 10], 20: [10, 30], 40: [30, 50], 60: [50, 70], 80: [70, 90], 100: [90, 110],
120: [110, 130], 140: [130, 150], 160: [150, 170]}
    histo_center = np.zeros((ip_img.shape[0], ip_img.shape[1]))

    # First the bin centers for each pixel are identified
    for i in range(ip_theta.shape[0]):
        for j in range(ip_theta.shape[1]):
            for x in bins_table.keys():
                if bins_table[x][0] <= ip_theta[i][j] < bins_table[x][1]:
                    histo_center[i][j] = x

    # hog histograms for each cell are created
    histo_list = np.zeros((int(ip_img.shape[0]/8), int(ip_img.shape[1]/8), 9))

    a = 0
    k = 0
    for x in range(int(ip_theta.shape[0]/8)):
        b = 0
        c = 0
        for y in range(int(ip_theta.shape[1]/8)):
            for i in range(a, a+8):
                for j in range(b, b+8):
                    bin_no = int(histo_center[i][j]/20)
                    if ip_theta[i][j] == histo_center[i][j]:
                        histo_list[k][c][bin_no] += ip_img[i][j]
```

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

```
else:
    diff1 = abs(bins_table[histo_center[i][j]][0] - ip_theta[i][j])
    diff2 = abs(bins_table[histo_center[i][j]][1] - ip_theta[i][j])
    if ip_theta[i][j] < histo_center[i][j]:
        histo_list[k][c][bin_no] += (diff2/20)*ip_img[i][j]
        histo_list[k][c][bin_no - 1] += (diff1/20)*ip_img[i][j]
    else:
        histo_list[k][c][bin_no] += (diff1/20)*ip_img[i][j]
    if bin_no == 8:
        histo_list[k][c][0] += (diff2/20)*ip_img[i][j]
    else:
        histo_list[k][c][bin_no + 1] += (diff2/20)*ip_img[i][j]

b += 8
c += 1
a += 8
k += 1

# block_array will store the normalized histogram block wise
block_array = np.zeros((histo_list.shape[0] - 1, histo_list.shape[1] - 1, 36))

# hog_descriptor will create a one dimensional array of all the normalized hirtograms
hog_descriptor_op = []
for i in range(histo_list.shape[0] - 1):
    for j in range(histo_list.shape[1] - 1):
        l1 = histo_list[i][j]
        l2 = histo_list[i+1][j]
        l3 = histo_list[i][j+1]
        l4 = histo_list[i+1][j+1]

        norm_factor = math.sqrt(np.sum(np.square(l1)) + np.sum(np.square(l2)) +
np.sum(np.square(l3)) + np.sum(np.square(l4)))
        np.seterr(divide='ignore', invalid='ignore')
        l1_, l2_, l3_, l4_ = np.true_divide(l1, norm_factor), np.true_divide(l2, norm_factor),
np.true_divide(l3, norm_factor), np.true_divide(l4, norm_factor)
        block_array[i][j] = np.concatenate((l1_, l2_, l3_, l4_))
        hog_descriptor_op.extend(l1_.tolist())
        hog_descriptor_op.extend(l2_.tolist())
        hog_descriptor_op.extend(l3_.tolist())
        hog_descriptor_op.extend(l4_.tolist())

return np.array(hog_descriptor_op).reshape((1, len(hog_descriptor_op)))
```

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

```
# main function
def hog(self):
    image_name = self.image_name
    img = imread(image_name)
    gs_img = self.greyscale_operation(img)
    grad_img, theta = self.gradient_operation(gs_img)
    hog_descriptor = self.hog_operation(grad_img, theta)
    return hog_descriptor
```

2. create\_data.py:

```
import HoG as HoG
import numpy as np
import os

# function to create data
def create_data():
    files = os.listdir('./positive_train')
    for i in range(len(files)):
        hog_obj = HoG.HOG('./positive_train/'+files[i])
        hog = hog_obj.hog()
        if i == 0:
            train_X_1 = hog
        else:
            train_X_1 = np.vstack((train_X_1, hog))

    train_y_1 = np.full((train_X_1.shape[0],1), 1)

    files = os.listdir('./negative_train')
    for i in range(len(files)):
        hog_obj = HoG.HOG('./negative_train/'+files[i])
        hog = hog_obj.hog()
        if i == 0:
            train_X_0 = hog
        else:
            train_X_0 = np.vstack((train_X_0, hog))

    train_y_0 = np.full((train_X_0.shape[0],1), 0)

    train_x = np.vstack((train_X_1, train_X_0))
    train_x = np.nan_to_num(train_x)
    train_y = np.vstack((train_y_1, train_y_0))
```

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

```
train = np.append(train_x, train_y, 1)

# saving training data to csv file
np.savetxt("train.csv",train,delimiter=",")

files = os.listdir('./positive_test')
for i in range(len(files)):
    hog_obj = HoG.HOG('./positive_test/'+files[i])
    hog = hog_obj.hog()
    if i == 0:
        test_X_1 = hog
    else:
        test_X_1 = np.vstack((test_X_1, hog))

test_y_1 = np.full((test_X_1.shape[0],1), 1)

files = os.listdir('./negative_test')
for i in range(len(files)):
    hog_obj = HoG.HOG('./negative_test/'+files[i])
    hog = hog_obj.hog()
    if i == 0:
        test_X_0 = hog
    else:
        test_X_0 = np.vstack((test_X_0, hog))

test_y_0 = np.full((test_X_0.shape[0],1), 0)

test_x = np.vstack((test_X_1, test_X_0))
test_x = np.nan_to_num(test_x)
test_y = np.vstack((test_y_1, test_y_0))
test = np.append(test_x, test_y, 1)

# saving testing data to csv file
np.savetxt("test.csv",test,delimiter=",")

create_data()
```

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

3. neural\_net.py

import numpy as np

class MLP:

```
def __init__(self, w1, b1, w2, b2, lr):  
    self.fc1 = FCLayer(w1, b1, lr)  
    self.rel = ReLU()  
    self.fc2 = FCLayer(w2, b2, lr)  
    self.sig = Sigmoid()
```

```
# function to calculate mean squared error  
def MSE(self, prediction, target):  
    return (0.5*(target-prediction)**2).sum()
```

```
# function to calculate the error  
def MSEGrad(self, prediction, target):  
    return -(target - prediction)
```

```
# training neural network  
def train(self, X, y, steps):  
    stop = False  
    prev_loss = 0.0  
    s = 0
```

very low

```
# training will end when either epoch iterations are completed or when error is
```

```
while stop != True and s != steps:  
    i = s % y.size  
    xi = np.expand_dims(X[i], axis=0)  
    yi = np.expand_dims(y[i], axis=0)  
  
    pred = self.fc1.forward(xi)  
    pred = self.rel.forward(pred)  
    pred = self.fc2.forward(pred)  
    pred = self.sig.forward(pred)  
    loss = self.MSE(pred, yi)  
  
    if round(abs(loss - prev_loss), 6) == 0.0:  
        print("Epochs:", s/y.size+1)  
        stop = True  
        break
```



Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

```
prev_loss = loss
grad = self.MSEGrad(pred, yi)
grad = self.sig.backward(grad)
grad = self.fc2.backward(grad)
grad = self.rel.backward(grad)
grad = self.fc1.backward(grad)
s += 1
```

```
# prediction using trained NN
def predict(self, X):
    pred = self.fc1.forward(X)
    pred = self.rel.forward(pred)
    pred = self.fc2.forward(pred)
    pred = self.sig.forward(pred)
    # pred = np.round(pred)
    return np.ravel(pred)
```

class FCLayer:

```
def __init__(self, w, b, lr):
    self.lr = lr
    self.w = w
    self.b = b

# forward pass
def forward(self, input):
    self.input = input
    h = np.dot(input, self.w) + self.b
    return h

# backward pass
def backward(self, gradients):
    input = self.input
    x_ = np.dot(gradients, self.w.T)
    self.w = self.w - np.dot(input.T, gradients)*self.lr
    self.b = self.b - gradients*self.lr
    return x_
```

class Sigmoid:

```
def __init__(self):
    None
```

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

```
def sigmoid_func(self, a):  
    return 1/(1+np.exp(-a))
```

#forward pass

```
def forward(self, input):  
    self.input = input  
    sig_val = self.sigmoid_func(input)  
    return sig_val
```

# backward pass

```
def backward(self, gradients):  
    input = self.input  
    sig_val_back = gradients*(1 - self.sigmoid_func(input))*self.sigmoid_func(input)  
    return sig_val_back
```

class ReLU:

```
    def __init__(self):  
        None
```

# forward pass

```
def forward(self, input):  
    self.input = input  
    input[input<0] = 0  
    return input
```

# backward pass

```
def backward(self, gradients):  
    input = self.input  
    input[input < 0] = 0  
    input[input > 0] = 1  
    input *= gradients  
    return input
```

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

4. run\_NN.py

```
import numpy as np
import neural_net as model

# function to load saved data
def load_data(path):
    data = np.genfromtxt(path, delimiter=',', dtype=float)
    return data[:, :-1], data[:, -1].astype(int)

train_x, train_y = load_data("train.csv")
test_x, test_y = load_data("test.csv")

# MLP Training
# learning rate
lr = 0.01

# random weight initialization
w1 = np.random.normal(0, .1, size=(train_x.shape[1], 250))
w2 = np.random.normal(0, .1, size=(250, 1))
b1 = np.random.normal(0, .1, size=(1, 250))
b2 = np.random.normal(0, .1, size=(1, 1))

mlp = model.MLP(w1, b1, w2, b2, lr)

# set epoch values
epoch = 400
steps = epoch * train_y.size

# training neural network
mlp.train(train_x, train_y, steps)

# evaluation function to calculate accuracy
def evaluate(solutions, real):
    if(solutions.shape != real.shape):
        raise ValueError("Output is wrong shape.")
    predictions = np.array(solutions)
    labels = np.array(real)
    return (predictions == labels).sum() / float(labels.size)

# predicting on test data
solutions = mlp.predict(test_x)
```

Project 2: Human Detection  
Computer Vision (CS 6643)  
Bhushan Manohar Newalkar (bmn258)

```
# printing NN output  
print("NN Output:")  
print(solutions)
```

```
# printing predictions  
print("Predictions")  
solutions = np.round(solutions)  
print(solutions)
```

```
# printing evaluation accuracy  
print(evaluate(solutions, test_y))
```